PDF Ingestion Note

PDF files often contain the official version of a policy, design note, or operating procedure. The ingestion service extracts the text layer page by page, records the page count, and passes the resulting text to the same chunking pipeline used for Markdown and plain-text files.

Scanned PDFs need optical character recognition before their text can be indexed. Tables and figures need specialized extractors when their layout carries meaning. Those extractors can emit Document records with the same source and page metadata used by the text pipeline.